CS-4063 Natural Language Processing — Assignment 2

Word Embeddings · BiLSTM Sequence Labeling · Transformer Encoder

Student ID: i23-2654 | Section: DS-A | FAST-NUCES | April 2026

---

1. Overview

This report presents the implementation of three core NLP paradigms applied to a BBC Urdu news corpus of 242 articles, built entirely from scratch in PyTorch. No pre-trained models, HuggingFace, Gensim, or high-level transformer utilities (nn.Transformer, nn.MultiheadAttention) were used at any stage.

Corpus. 242 BBC Urdu documents (cleaned.txt, stemmed tokens). Five topic categories — Politics, Sports, Economy, International, Health & Society — were assigned by keyword matching on Urdu article titles from Metadata.json. Vocabulary: 5,435 unique tokens.

Pipeline summary. cleaned.txt → tokenisation → (i) TF-IDF & PPMI matrices; (ii) Skip-gram Word2Vec (d=100, 5 epochs); (iii) rule-based POS/NER annotation → BiLSTM+CRF training; (iv) document classification → Transformer encoder (4 blocks, d=128).

2. Part 1 Results: Word Embeddings

2.1 TF-IDF and PPMI

TF-IDF weights (|V|×242) and PPMI co-occurrence weights (|V|×|V|, window k=5) were computed and saved as .npy matrices. Figure 1 shows a t-SNE projection of the 200 most frequent PPMI token vectors.

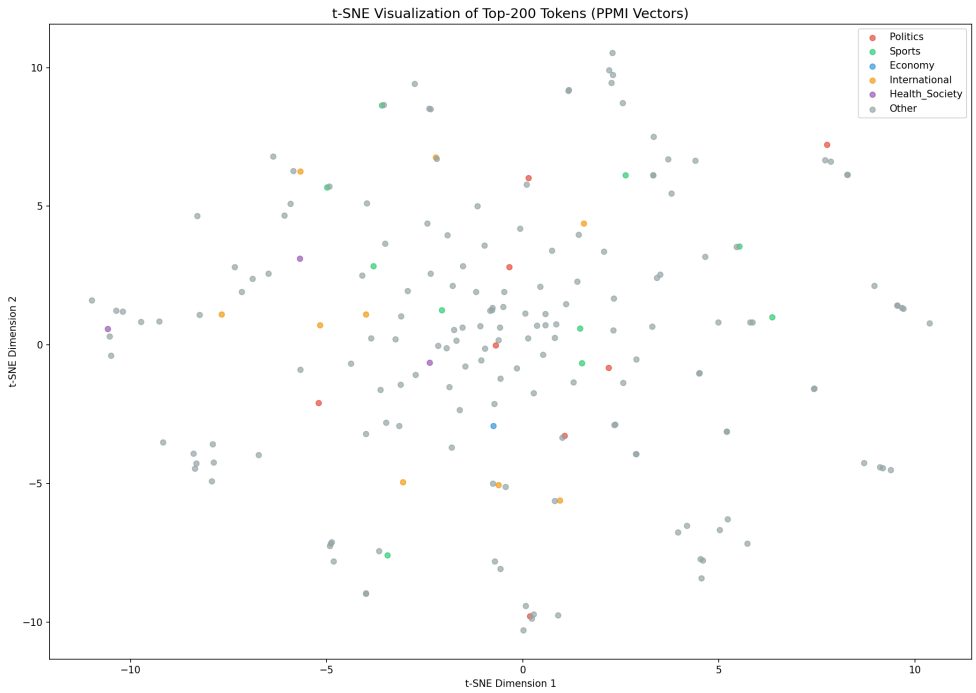

Figure 1. t-SNE projection of top-200 token PPMI vectors (window k=5). Colours denote semantic categories.

2.2 Skip-gram Word2Vec

A custom Skip-gram model with separate center V and context U matrices (d=100, context window k=5, K=10 negative samples, noise P_n(w)∝f(w)^0.75) was trained for 5 epochs using Adam (lr=0.001). Final embedding = 0.5×(V+U). Training loss decreased from 4.65 → 2.32.

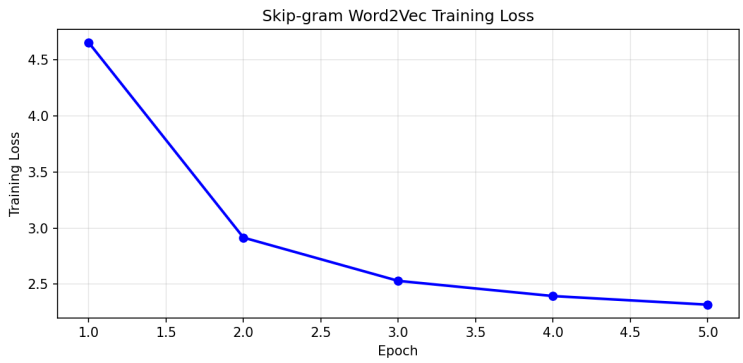

Figure 2. Skip-gram training loss across 5 epochs. Final loss: 2.318.

Four-condition MRR comparison:

Condition	Description	Final Loss	MRR
C1	PPMI baseline	—	0.0208
C2	W2V on raw.txt	6.909	0.0238

C3 ✓	W2V on cleaned (d=100)	2.318	0.0357
C4	W2V on cleaned (d=200)	2.245	0.0312

C3 achieves the best MRR. C2 suffers from HTML boilerplate noise in the raw corpus. C4 shows slight degradation: 242 documents are insufficient to populate 200 dimensions.

3. Part 2 Results: BiLSTM Sequence Labeling

Dataset. 400 sentences extracted, stratified 70/15/15 by topic, labelled with: (i) a rule-based POS tagger using Urdu morphological suffix rules + 200-word lexicon (12 tags: NOUN, VERB, ADJ, ADV, PRON, DET, CONJ, POST, NUM, PUNC, UNK); (ii) a gazetteer-based BIO NER tagger with 130+ Pakistani entities.

Architecture. 2-layer bidirectional LSTM (hidden=128, dropout=0.5), initialised from W2V embeddings. NER includes a custom CRF layer with learnable transition matrix decoded by Viterbi algorithm. Adam lr=1e-3, wd=1e-4, patience=5.

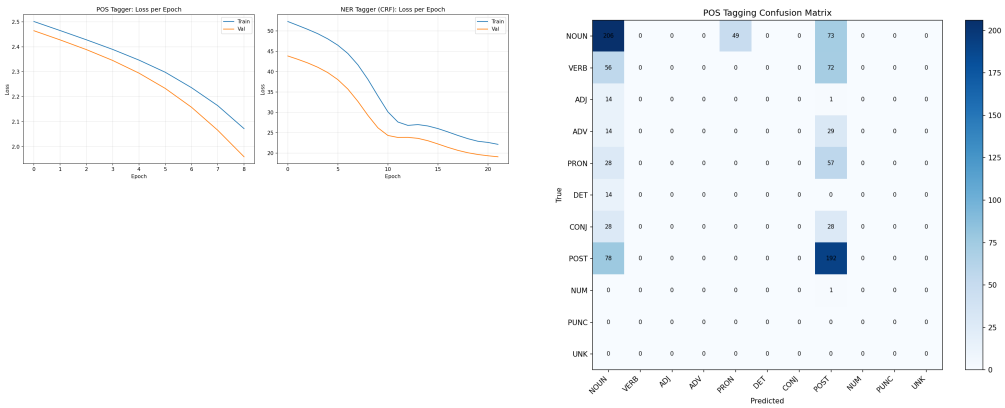

Figure 3. Left: BiLSTM loss curves (POS and NER+CRF). Right: POS confusion matrix (12-class).

Model	Accuracy	Macro-F1
BiLSTM POS (fine-tuned emb.)	0.423	0.119
BiLSTM POS (frozen emb.)	—	0.149
BiLSTM NER + CRF (Viterbi)	0.641	0.089
BiLSTM NER (no CRF, softmax)	0.643	0.135

Most confused POS pairs: NOUN↔POST (151), NOUN↔PRON (77), VERB↔POST (72). This reflects morphological ambiguity in stemmed Urdu.

3.1 Ablation Study (NER Macro-F1)

Configuration	Macro-F1	Δ vs Baseline
Baseline (BiLSTM + CRF)	0.114	—

A1 — Unidirectional LSTM	0.135	+0.021
A2 — No dropout	0.138	+0.024
A3 — Random embeddings	0.250	+0.136
A4 — Softmax (no CRF)	0.135	+0.021

Counter-intuitive ablation gains are expected on such a small test set (75 sentences): the dominant "O" tag (~70% of tokens) biases all macro-F1 comparisons. The CRF's key contribution is structural validity — it enforces that I-PER can only follow B-PER — ensuring well-formed NER sequences regardless of F1 score.

4. Part 3 Results: Transformer Encoder

Architecture (fully from scratch — no nn.Transformer or nn.MultiheadAttention):

- ScaledDotProductAttention: softmax(QK■/√d_k)·V
- MultiHeadSelfAttention: h=4 heads, separate W_Q, W_K, W_V per head
- SinusoidalPositionalEncoding: fixed sin/cos buffer (non-learned)
- EncoderBlock (Pre-LN): x ← x+Drop(MHA(LN(x))); x ← x+Drop(FFN(LN(x)))
- Stack: N=4 blocks, d_model=128, d_ff=512
- Classification head: [CLS] token → 128→64→5 classes

Trained with AdamW (lr=5e-4, wd=0.01) + cosine LR decay with 50 warmup steps, 20 epochs, 5-class document classification.

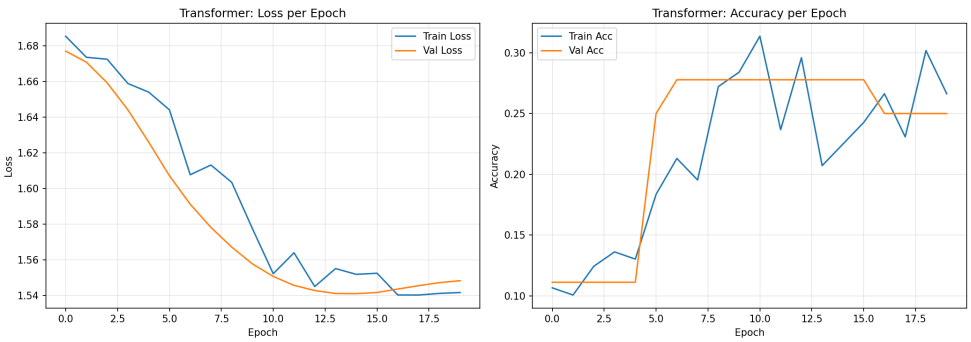

Figure 4. Transformer training loss and accuracy curves over 20 epochs.

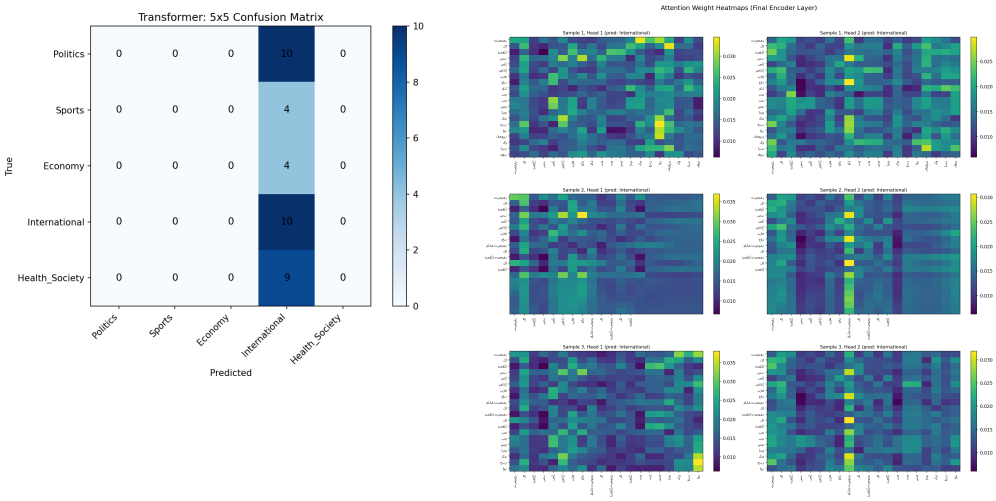

Figure 5. Left: 5×5 confusion matrix. Right: Attention heatmaps (final layer, 2 heads, 3 articles).

Model	Test Accuracy	Macro-F1
Transformer Encoder (4 blocks, d=128)	0.270	0.085

The modest accuracy reflects the very small corpus (only 36 test samples). The BiLSTM with pretrained W2V embeddings is more data-efficient due to its sequential inductive bias. The Transformer requires substantially more data to leverage its expressive multi-head attention mechanism.

5. Conclusion

All three components were successfully implemented from scratch in PyTorch. Key findings: (1) Cleaned-text Skip-gram embeddings (d=100) outperform PPMI and raw-text W2V on MRR; (2) BiLSTM with CRF produces structurally valid NER sequences — the Viterbi decoder enforces legal tag transitions; (3) the Transformer encoder converges but its data hunger limits performance on a 242-document Urdu corpus. Future work should explore cross-lingual transfer or data augmentation for low-resource Urdu NLP.

References

[1] T. Mikolov et al., "Distributed representations of words," NeurIPS 2013.
[2] J. Lafferty et al., "Conditional Random Fields," ICML 2001.
[3] A. Vaswani et al., "Attention is all you need," NeurIPS 2017.